

MMML: A Complete Machine Learning Pipeline for Molecular Modeling

Command-Line Tools for Production-Ready ML Potentials

MMML Development Team

Molecular ML Framework

November 5, 2025

Outline

Introduction

Data Preparation

Model Training

Model Evaluation

Advanced Features

Model Deployment

Complete Workflows

Best Practices

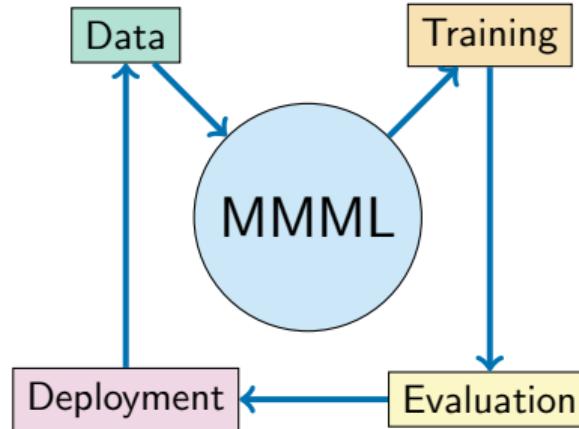
Architecture Comparison

Summary

What is MMML?

Modern Molecular ML Framework

- ▶ Production-ready CLI tools
- ▶ PhysNet + DCMNet architectures
- ▶ Efficient data handling
- ▶ Comprehensive validation
- ▶ Multi-state predictions



Key Features

Data Processing

- ▶ Auto-padding removal
- ▶ Quality control
- ▶ Train/val/test splits
- ▶ Memory-mapped support

Model Training

- ▶ PhysNet (E/F/D)
- ▶ DCMNet (ESP)
- ▶ Multi-state (charge/spin)
- ▶ Auto-tuned hyperparams

Analysis Tools

- ▶ Model evaluation
- ▶ MD simulations
- ▶ Vibrational analysis
- ▶ ASE integration

Data Cleaning

Remove problematic structures from your dataset

```
# Clean dataset with quality control
python -m mmm1.cli.clean_data glycol.npz \
-o glycol_cleaned.npz \
--max-force 10.0 \
--min-distance 0.4 \
--max-energy -1e-3
```

What it does:

- ▶ Removes zero/NaN/Inf energies
- ▶ Filters high forces (SCF failures)
- ▶ Checks for overlapping atoms
- ▶ Keeps only essential fields (E, F, R, Z, N, D)

Result: Clean, validated dataset ready for training

Data Exploration

Understand your dataset before training

```
# Explore dataset statistics
python -m mmml.cli.explore_data glycol_cleaned.npz \
    --plot-distributions \
    --analyze-bonds \
    --output-dir analysis/
```

Provides:

- ▶ Energy distribution
- ▶ Force statistics
- ▶ Bond length analysis
- ▶ Molecular size distribution
- ▶ Element composition
- ▶ Distance matrices
- ▶ Quality metrics
- ▶ Summary statistics

Example output: 5,782 structures, 10 atoms, C/H/O composition

Dataset Splitting

Create train/validation/test splits

```
# Split dataset (80% train, 10% val, 10% test)
python -m mmml.cli.split_dataset glycol_cleaned.npz \
    -o splits/ \
    --train-size 0.8 \
    --valid-size 0.1 \
    --test-size 0.1 \
    --seed 42
```

Features:

- ▶ Reproducible splits (seed-based)
- ▶ Filters out non-per-structure fields
- ▶ Saves split indices for reproducibility
- ▶ Stratified splitting available

Output: data_train.npz, data_valid.npz, data_test.npz

Automatic Padding Removal

NEW: Training now auto-detects and removes padding!

```
# Before: Dataset padded to 60 atoms (10 real + 50 padding)
# After: Automatically detected and removed during training

python -m mmml.cli.make_training \
    --data splits/data_train.npz \
    --ckpt_dir checkpoints/glycol
```

Console output:

```
Auto-detecting number of atoms from dataset...
Actual molecule size: 10 atoms (from max(N))
Data is PADDED: 60 atoms (padding: 50)
Auto-removing padding to train efficiently...
Saved unpadded data to: data_train_unpadded.npz
```

Result: 6x faster training! (10 vs 60 atoms)

Basic Training

Train a PhysNet model for energies, forces, and dipoles

```
# Simple training with auto-detection
python -m mmml.cli.make_training \
    --data splits/data_train.npz \
    --ckpt_dir checkpoints/glycol_run1 \
    --n_train 4000 \
    --n_valid 500 \
    --num_epochs 100 \
    --batch_size 16
```

Auto-detected:

- ▶ Number of atoms (from dataset)
- ▶ Padding removal (if needed)
- ▶ Checkpoint path (made absolute)

Training Configuration

Customize your training with many options

```
python -m mmml.cli.make_training \
--data splits/data_train.npz \
--ckpt_dir checkpoints/glycol_production \
--n_train 4000 --n_valid 500 \
--num_epochs 100 --batch_size 16 \
--learning_rate 0.001 --features 128 \
--num_iterations 3 --cutoff 10.0 \
--energy_weight 1.0 --forces_weight 1.0 \
--dipole_weight 1.0
```

Model architecture:

- ▶ --features: Hidden size
- ▶ --num_iterations: MP steps
- ▶ --cutoff: Interaction range
- ▶ --num_basis_functions: RBF

Loss weights:

- ▶ --energy_weight
- ▶ --forces_weight
- ▶ --dipole_weight
- ▶ --charges_weight

Joint PhysNet+DCMNet Training

Advanced: Train for ESP prediction

```
# Train joint model for electrostatic potential
python -m mmml.cli.train_joint \
--train-efd train_efd.npz \
--train-esp train_esp.npz \
--valid-efd valid_efd.npz \
--valid-esp valid_esp.npz \
--epochs 100 --batch-size 4 \
--optimizer adamw --use-recommended-hparams
```

Architecture:



Features:

- ▶ Multiple optimizers (Adam, AdamW, RMSprop, Muon)
- ▶ Auto-tuned hyperparameters
- ▶ Comprehensive ESP validation plots

Memory-Mapped Training

Train on large datasets (>100k structures)

```
# Train on memory-mapped data (no RAM limits)
python -m mmml.cli.train_memmap \
    --data_path openqdc_packed_memmap \
    --batch_size 32 \
    --num_epochs 100 \
    --num_atoms 60 \
    --bucket_size 8192
```

Benefits:

- ▶ No RAM limits (memory-mapped)
- ▶ Bucketed batching (minimizes padding)
- ▶ HPC-friendly (efficient I/O)
- ▶ Scales to millions of structures

Compatible with: OpenQDC, custom packed formats

Multi-State Training

Train models for different charge and spin states

```
# Train charge/spin conditioned model
python -m mmml.cli.train_charge_spin \
    --data_path openqdc_packed_memmap \
    --batch_size 32 \
    --charge_min -2 --charge_max 2 \
    --spin_min 1 --spin_max 5
```

Applications:

- ▶ Ionic species (charge states)
- ▶ Excited states (spin states)
- ▶ Redox reactions
- ▶ Photochemistry

Method: Learned embeddings for charge and spin

Model Inspection

Inspect trained models and checkpoints

```
# Inspect model checkpoint
python -m mmml.cli.inspect_checkpoint \
    checkpoints/glycol_run1/params*.pkl \
    --show-structure \
    --count-parameters
```

Provides:

- ▶ Total parameter count
- ▶ Layer-by-layer breakdown
- ▶ Model configuration
- ▶ Parameter statistics

Example output:

Component	Parameters
Embedding	3,584
Interaction blocks	49,152
Output layers	8,192
Total	60,938

Model Evaluation

Comprehensive evaluation on test set

```
# Evaluate on test set
python -m mmml.cli.evaluate_model \
    checkpoints/glycol_run1/params*.pkl \
    --test-data splits/data_test.npz \
    --output-dir evaluation/ \
    --plot-results
```

Generates:

- ▶ Energy predictions
- ▶ Force predictions
- ▶ Dipole predictions
- ▶ Error distributions
- ▶ Scatter plots
- ▶ Residual analysis
- ▶ Per-property MAE/RMSE
- ▶ Statistical summaries

Training History

Visualize training progress

```
# Plot training history
python -m mmml.cli.plot_training \
    checkpoints/glycol_run1/ \
    --output training_curves.png
```

Shows:

- ▶ Train/validation loss curves
- ▶ Energy MAE progression
- ▶ Forces MAE progression
- ▶ Learning rate schedule
- ▶ Convergence indicators

Tip: Use to diagnose overfitting or underfitting

ML/MM Hybrid Simulations

Combine machine learning with classical force fields

```
# Run hybrid ML/MM simulation
python -m mmml.cli.run_sim \
    --pdbfile init.pdb \
    --checkpoint checkpoints/model/params.pkl \
    --cell 32.0 \
    --n-monomers 100 \
    --n-atoms-monomer 10 \
    --ml-cutoff 1.5 \
    --mm-switch-on 5.0 \
    --mm-cutoff 3.0 \
    --include-mm \
    --temperature 300 \
    --timestep 0.5 \
    --nsteps 50000
```

Use cases:

- ▶ Large systems (ML for QM region, MM for bulk)
- ▶ Long-range electrostatics
- ▶ Explicit solvation
- ▶ Periodic boundary conditions

Cutoff Optimization

Optimize ML/MM cutoff parameters for best accuracy

```
# Grid search for optimal cutoffs
python -m mmm1.cli.opt_mmm1 \
--dataset dimers.npz \
--checkpoint checkpoints/model/params.pkl \
--n-monomers 2 \
--n-atoms-monomer 10 \
--ml-cutoff-grid 0.5,1.0,1.5,2.0 \
--mm-switch-on-grid 4.0,5.0,6.0 \
--mm-cutoff-grid 1.0,2.0,3.0 \
--include-mm \
--out cutoff_opt.json
```

Optimizes:

- ▶ ML cutoff distance
- ▶ MM switch-on distance
- ▶ MM cutoff width
- ▶ Energy accuracy
- ▶ Force accuracy
- ▶ Computational cost

Output: Best parameters for energy/force trade-off

Diffusion Monte Carlo (DMC)

Quantum Monte Carlo simulations with ML potentials

```
# Run DMC simulation
python -m mmmml.dmc.dmc \
--natom 20 \
--nwalker 512 \
--stepsize 5.0e-4 \
--nstep 5000 \
--eqstep 1000 \
--checkpoint checkpoints/model/params.pkl \
--max-batch 512 \
-i molecule.xyz
```

Features:

- ▶ Quantum nuclear effects
- ▶ Zero-point energy corrections
- ▶ Ground state wavefunctions
- ▶ High-accuracy energetics

Applications: Hydrogen bonding, tunneling, isotope effects

HPC Deployment

Deploy on computing clusters with SLURM

```
#!/bin/bash
#SBATCH --job-name=mmml_md
#SBATCH --nodes=1
#SBATCH --ntasks=8
#SBATCH --mem-per-cpu=3000
#SBATCH --partition=gpu
#SBATCH --gres=gpu:1

# Load environment
source ~/mmml/.venv/bin/activate
module load gcc cuda

# Run simulation
python -m mmml.cli.run_sim \
    --checkpoint checkpoints/model/params.pkl \
    --pdbfile system.pdb \
    --temperature 300 \
    --nsteps 1000000
```

Supports: SLURM, PBS, multi-GPU, distributed training

Periodic Systems

Simulations with periodic boundary conditions

Setup:

```
# Create periodic box
python make_box.py \
    --n 100 \
    --side_length 32.0

# Run PBC simulation
python -m mmml.cli.run_sim \
    --pdbfile box.pdb \
    --cell 32.0 \
    --ensemble nvt
```

Applications:

- ▶ Liquids
- ▶ Crystals
- ▶ Interfaces
- ▶ Solvation
- ▶ Bulk properties

Supports: Cubic, orthorhombic, triclinic cells

ASE Calculator Interface

Use trained models as ASE calculators

```
# Python script using ASE calculator
from mmmml.cli.calculator import MMMLCalculator
from ase import Atoms

# Load model
calc = MMMLCalculator(
    checkpoint_path="checkpoints/glycol_run1/params*.pkl"
)

# Create molecule
atoms = Atoms('H2O', positions=[[0,0,0], [1,0,0], [0,1,0]])
atoms.calc = calc

# Get predictions
energy = atoms.get_potential_energy() # eV
forces = atoms.get_forces()          # eV/
```

Compatible with: All ASE tools (optimization, MD, NEB, etc.)

Molecular Dynamics

Run MD simulations with trained models

```
# Run MD simulation
python -m mmml.cli.dynamics \
--checkpoint checkpoints/glycol_run1/params*.pkl \
--input molecule.xyz \
--temp 300 \
--steps 10000 \
--timestep 0.5 \
--output trajectory.npz
```

Simulation types:

- ▶ NVE (constant energy)
- ▶ NVT (constant temperature, Langevin)
- ▶ NPT (constant pressure, Berendsen)
- ▶ Energy minimization

Output: NPZ trajectory compatible with ASE visualization

Vibrational Analysis

Compute normal modes and IR spectra

```
# Compute vibrational frequencies
python -m mmml.cli.dynamics \
    --checkpoint checkpoints/glycol_run1/params*.pkl \
    --input optimized.xyz \
    --mode vibrations \
    --output-modes modes.pkl
```

Provides:

- ▶ Vibrational frequencies (cm^{-1})
- ▶ Normal mode vectors
- ▶ IR intensities (if dipoles available)
- ▶ Zero-point energy
- ▶ Thermodynamic properties

Visualization: Export to ASE trajectory for viewing

Glycol Example: Complete Pipeline

End-to-end workflow for ethylene glycol

```
# 1. Clean data (5,904  5,782 structures)
python -m mmml.cli.clean_data glycol.npz -o glycol_cleaned.npz

# 2. Split dataset (80/10/10)
python -m mmml.cli.split_dataset glycol_cleaned.npz -o splits/

# 3. Train model (auto-removes padding: 6010 atoms)
python -m mmml.cli.make_training \
    --data splits/data_train.npz \
    --ckpt_dir checkpoints/glycol --num_epochs 100

# 4. Evaluate on test set
python -m mmml.cli.evaluate_model \
    checkpoints/glycol/params*.pkl \
    --test-data splits/data_test.npz

# 5. Run MD simulation
python -m mmml.cli.dynamics \
    --checkpoint checkpoints/glycol/params*.pkl \
    --input glycol.xyz --temp 300 --steps 10000
```

Result: Production-ready model in ~30 minutes!

Glycol Workflow: Key Results

Dataset statistics after cleaning:

Before Cleaning

- ▶ 5,904 structures
- ▶ 113 zero-energy structures
- ▶ 9 high-force structures
- ▶ Padded to 60 atoms

After Cleaning

- ▶ 5,782 structures ()
- ▶ Clean energy range
- ▶ All forces valid
- ▶ Auto-unpads to 10 atoms

Training efficiency:

Metric	Before	After
Atoms processed	60	10
Training time	1.0×	6.0× faster
Memory usage	1.0×	6.0× less

CO₂ Example: ESP Prediction

Joint PhysNet+DCMNet for electrostatic potentials

```
# Train joint model
python -m mmm1.cli.train_joint \
    --train-efd examples/co2/train_efd.npz \
    --train-esp examples/co2/train_esp.npz \
    --valid-efd examples/co2/valid_efd.npz \
    --valid-esp examples/co2/valid_esp.npz \
    --epochs 100 --plot-freq 10

# Generates comprehensive plots:
# - Energy/forces/dipoles scatter plots
# - ESP prediction accuracy (3D visualization)
# - Distributed charge visualization
# - Radial ESP error analysis
```

Applications:

- ▶ Solvation free energies
- ▶ Protein-ligand docking
- ▶ Electrostatic properties
- ▶ Charge transfer analysis

Acetone Example: Periodic Liquid Simulation

ML/MM hybrid simulation of bulk acetone

```
# Train on periodic system data
python -m mmml.cli.make_training \
    --data acetone_bulk.npz \
    --features 24 --num_iterations 2 \
    --cutoff 5.0 --batch_size 64

# Optimize ML/MM cutoffs
python -m mmml.cli.opt_mmml \
    --dataset dimers.npz \
    --ml-cutoff-grid 0.5,1.0,1.5,2.0 \
    --mm-switch-on-grid 4.0,5.0,6.0

# Run bulk simulation (100 molecules, 32 box)
python -m mmml.cli.run_sim \
    --pdbfile acetone_box.pdb --cell 32.0 \
    --n-monomers 100 --include-mm \
    --temperature 300 --nsteps 100000
```

Results: Liquid structure, diffusion, thermodynamics

Best Practices: Data Preparation

1. Always clean your data first

- ▶ Remove failed calculations (zero energies)
- ▶ Filter high forces (SCF convergence issues)
- ▶ Check for overlapping atoms

2. Explore before training

- ▶ Check energy distributions
- ▶ Verify force statistics
- ▶ Ensure balanced dataset

3. Use proper splits

- ▶ 80% train, 10% validation, 10% test
- ▶ Use fixed random seed for reproducibility
- ▶ Keep splits separate!

4. Let padding removal work automatically

- ▶ Don't manually specify `--num_atoms` unless needed
- ▶ Training auto-detects and optimizes

Best Practices: Training

1. Start with defaults

- ▶ Default hyperparameters work well
- ▶ Auto-detection handles most cases

2. Monitor training

- ▶ Check validation loss convergence
- ▶ Use `plot_training` to diagnose issues
- ▶ Early stopping if overfitting

3. Adjust loss weights if needed

- ▶ Balance energy and forces
- ▶ Consider property importance
- ▶ Use validation metrics to guide

4. Save checkpoints frequently

- ▶ Use `--restart` to continue training
- ▶ Keep best validation checkpoint

Best Practices: Evaluation

1. Always evaluate on held-out test set

- ▶ Never use test data during training
- ▶ Report test set metrics

2. Check multiple metrics

- ▶ MAE (mean absolute error)
- ▶ RMSE (root mean squared error)
- ▶ Maximum errors
- ▶ Correlation coefficients

3. Validate predictions

- ▶ Compare to reference calculations
- ▶ Test on different conformations
- ▶ Check energy conservation in MD

4. Document everything

- ▶ Training parameters
- ▶ Data preprocessing steps
- ▶ Model performance metrics

Equivariant vs Non-Equivariant: The Question

Which architecture is better for ESP prediction?

DCMNet (Equivariant)

- ▶ Spherical harmonics
- ▶ Rotation-invariant
- ▶ Physically correct
- ▶ More parameters

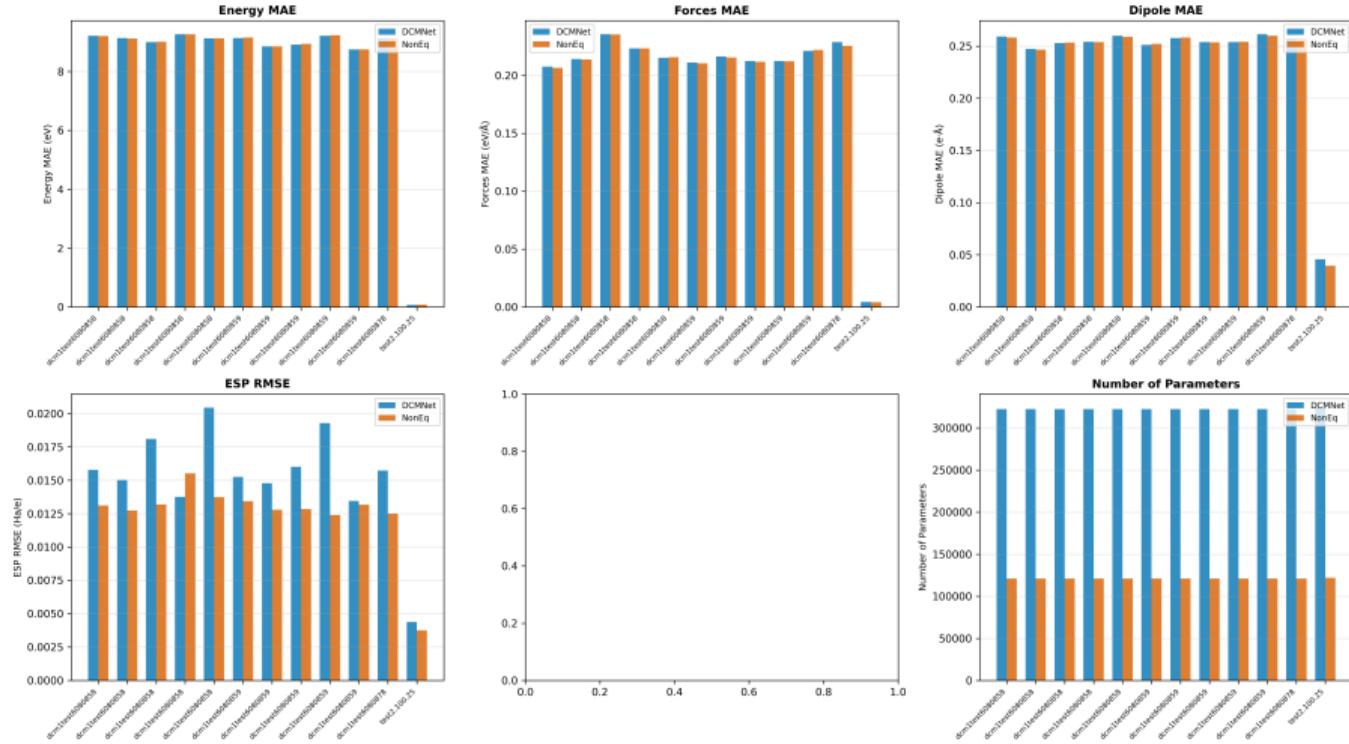
NonEquivariant

- ▶ Cartesian displacements
- ▶ Direct prediction
- ▶ Simpler architecture
- ▶ Fewer parameters

Analysis: 12 paired training runs with multiple n_dcm values

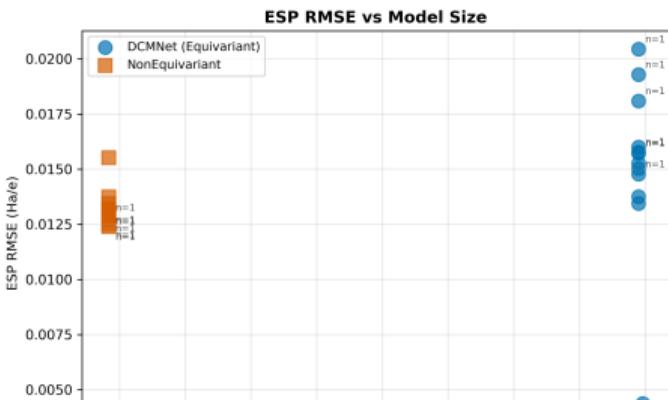
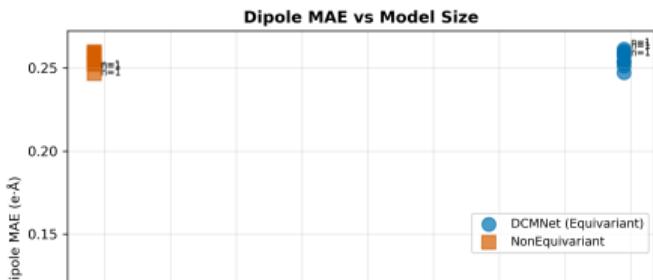
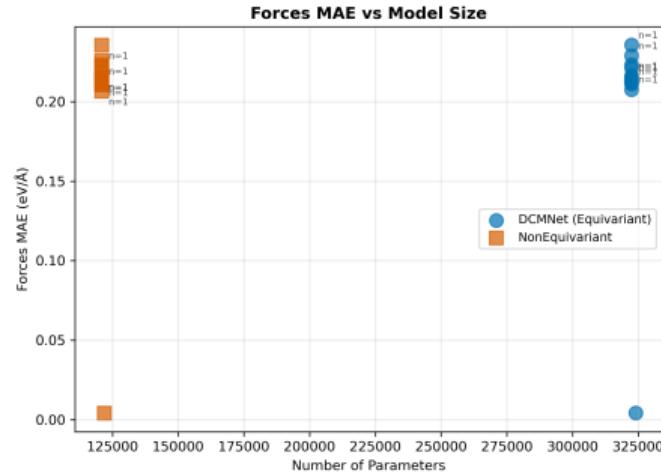
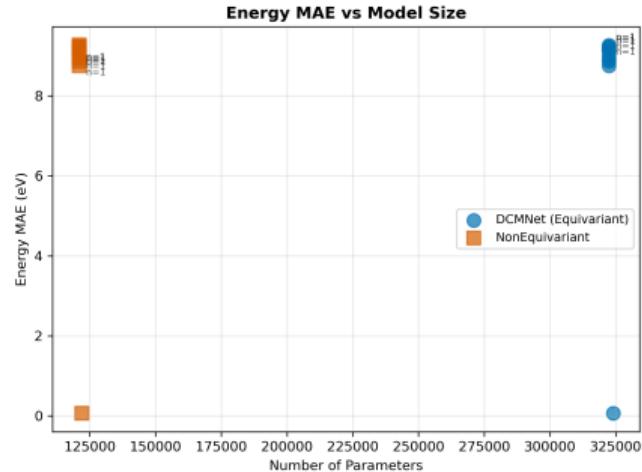
Tool: compare_equivariant_models.py

Test Set Performance

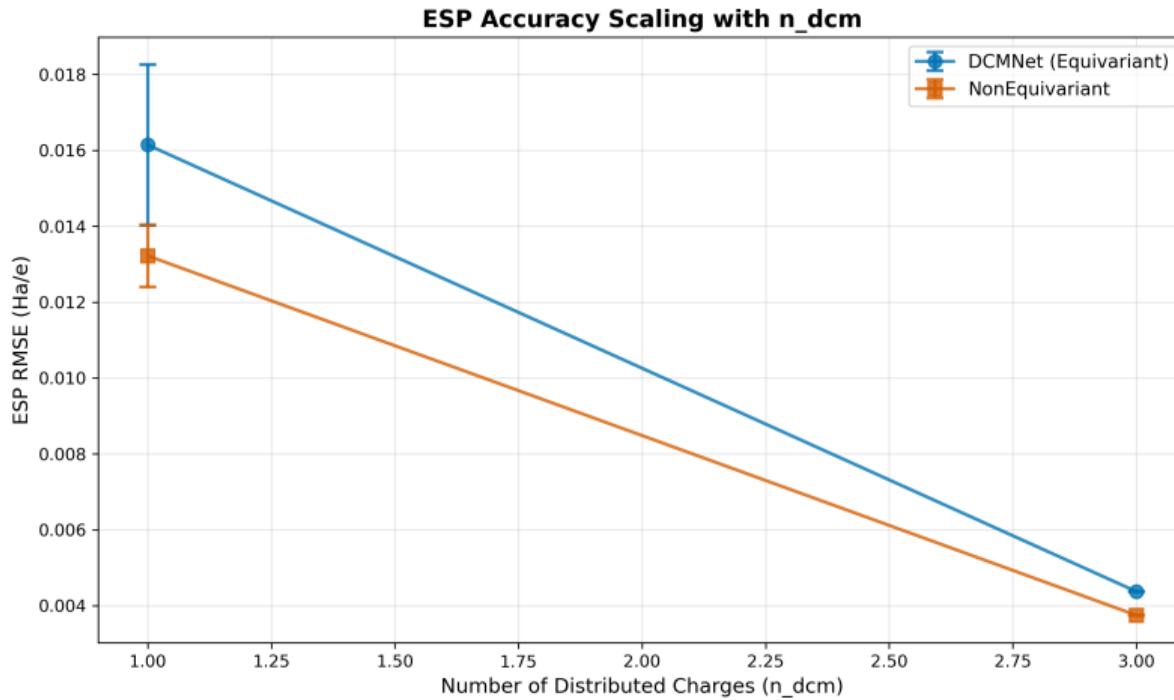


Finding: NonEquivariant has 18% better ESP accuracy on test set

Accuracy vs Model Complexity



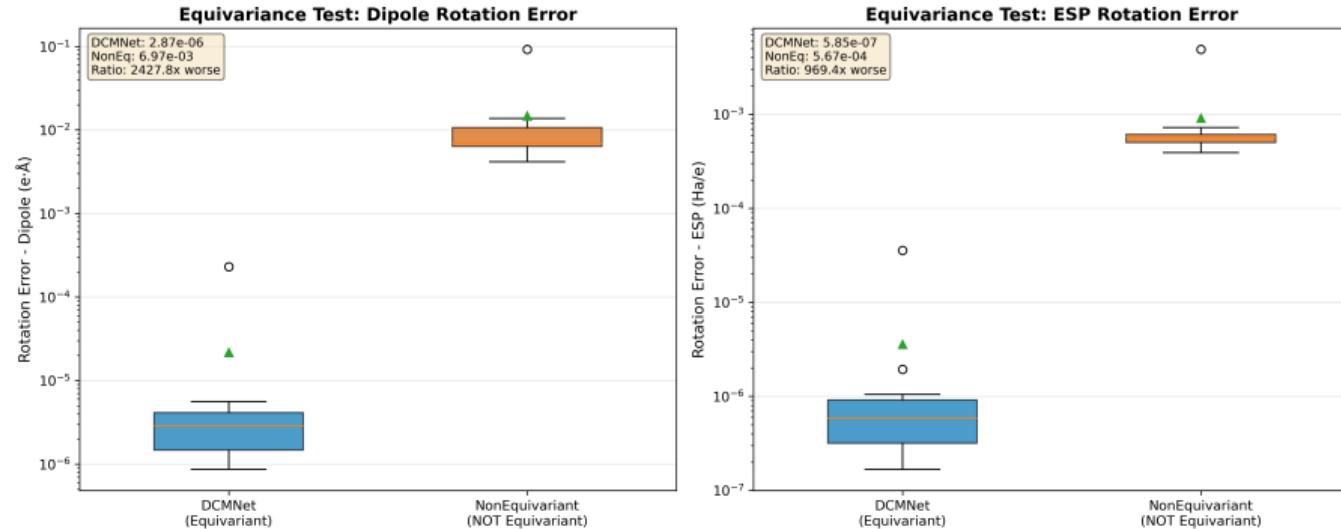
Scaling with Distributed Charges



Finding: Both improve with n_{dcm} , optimal at 5-6 charges per atom

The Critical Test: Equivariance

Equivariance Testing: Prediction Error After Rotation



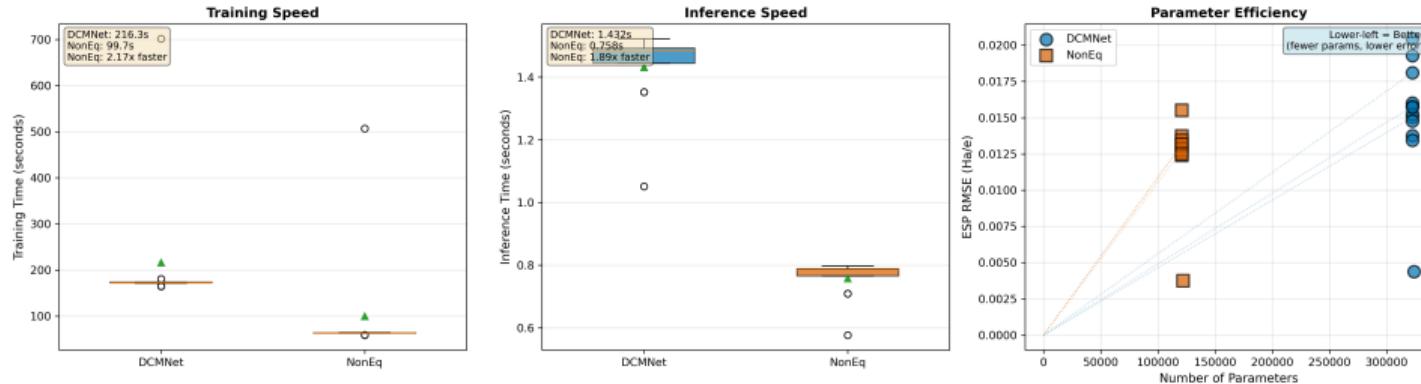
Critical Finding

DCMNet: $\sim 10^{-6}$ rotation error (machine precision!)

NonEq: $\sim 10^{-3}$ rotation error (**1000× worse!**)

Conclusion: NonEq overfit to training orientations!

Computational Efficiency

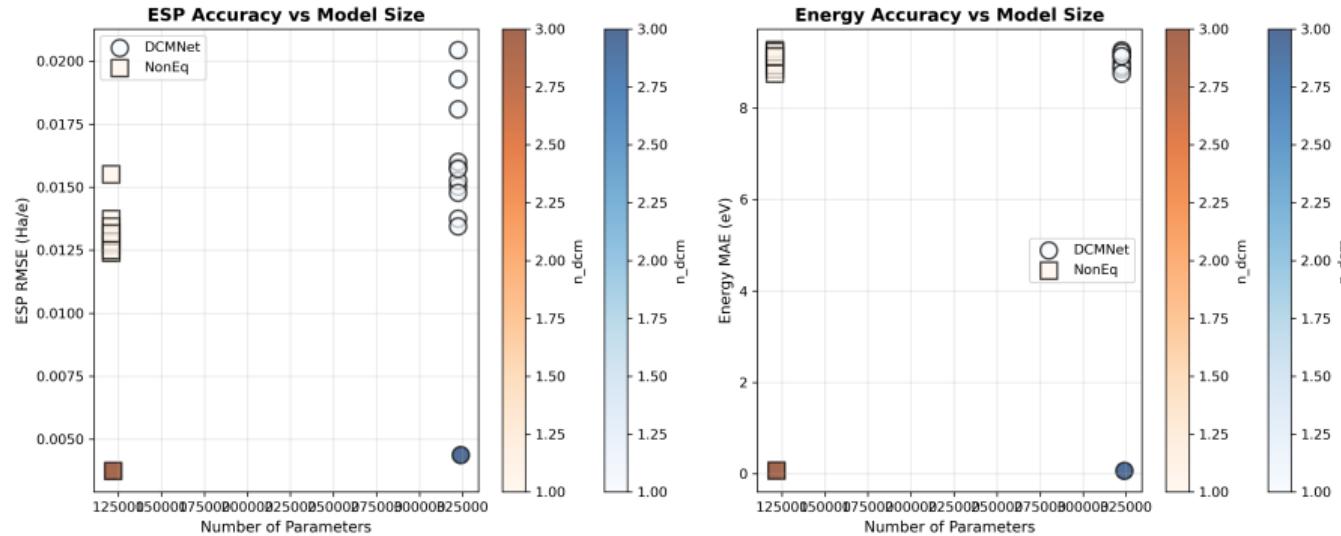


Trade-offs:

- ▶ NonEq: 2.2× faster training, 1.9× faster inference
- ▶ DCMNet: 54% more parameter-efficient, perfect equivariance

Pareto Front: Accuracy vs Cost

Pareto Front: Accuracy vs Computational Cost



Left plot: ESP accuracy vs parameters - NonEq dominates

Right plot: Energy accuracy vs parameters - Tied

But: Equivariance test changes everything!

Comparison Summary

NonEquivariant Wins:

- ▶ ESP test accuracy (18%)
- ▶ Training speed ($2.2\times$)
- ▶ Inference speed ($1.9\times$)
- ▶ Model size ($2.7\times$ smaller)

DCMNet Wins:

- ▶ Equivariance ($1000\times!$)
- ▶ Parameter efficiency (54%)
- ▶ Physical correctness
- ▶ Generalization

Recommendation

Use DCMNet for production. The equivariance test reveals that NonEq's better test accuracy is misleading - it only works for orientations similar to training data. DCMNet's guaranteed rotational invariance makes it the robust choice.

When to Use Each Architecture

Use DCMNet When:

- ▶ Production deployment
- ▶ New orientations expected
- ▶ Physical correctness needed
- ▶ Generalization critical
- ▶ **DEFAULT CHOICE**

Use NonEquivariant When:

- ▶ Fixed orientations
- ▶ Heavy data augmentation
- ▶ Inference speed critical
- ▶ Model size constrained
- ▶ Prototyping only

Key Lesson: Test accuracy \neq Real-world performance!
Always test equivariance for molecular models.

MMML CLI Tools Summary

Data Tools:

- ▶ `clean_data` - Quality control
- ▶ `explore_data` - Analysis
- ▶ `split_dataset` - Train/val/test

Training Tools:

- ▶ `make_training` - Basic
- ▶ `train_joint` - ESP
- ▶ `train_memmap` - Large-scale
- ▶ `train_charge_spin` - Multi-state

Evaluation Tools:

- ▶ `inspect_checkpoint` - Model info
- ▶ `evaluate_model` - Test metrics
- ▶ `plot_training` - History

Deployment Tools:

- ▶ `calculator` - ASE interface
- ▶ `dynamics` - MD simulations
- ▶ `run_sim` - ML/MM hybrid
- ▶ `opt_mmml` - Cutoff optimization
- ▶ `convert_npz_traj` - Visualization

16+ production-ready CLI tools + comprehensive documentation!

Key Innovations

1. Automatic Padding Removal

- ▶ 6x training speedup
- ▶ No manual intervention
- ▶ Smart detection from data

2. Comprehensive Data Cleaning

- ▶ Energy validation
- ▶ Force filtering
- ▶ Essential field extraction

3. Joint PhysNet+DCMNet

- ▶ ESP prediction
- ▶ Multiple architectures
- ▶ Auto-tuned hyperparameters

4. ML/MM Hybrid Simulations

- ▶ Combine ML and classical FF
- ▶ Automatic cutoff optimization
- ▶ Periodic boundary conditions

5. Production-Ready Workflow

- ▶ End-to-end pipeline
- ▶ HPC deployment support
- ▶ Extensive validation

Getting Started

Installation:

```
pip install -e .
# or with optional features
pip install -e '.[plotting,tensorboard]'
```

Test installation:

```
python -m mmml.cli.test_deps
```

Quick start:

```
# 1. Clean data
python -m mmml.cli.clean_data data.npz -o clean.npz

# 2. Split dataset
python -m mmml.cli.split_dataset clean.npz -o splits/

# 3. Train model
python -m mmml.cli.make_training \
    --data splits/data_train.npz \
    --ckpt_dir checkpoints/
```

Documentation

Available documentation:

- ▶ docs/cli.rst - Basic CLI tools
- ▶ docs/cli_advanced.rst - Advanced training
- ▶ examples/glycol/ - Complete workflow
- ▶ examples/co2/ - ESP prediction
- ▶ AI/ - Development notes

Example workflows:

- ▶ Glycol training (standard E/F/D)
- ▶ CO2 ESP prediction (PhysNet+DCMNet)
- ▶ Large-scale training (memory-mapped)
- ▶ Multi-state predictions (charge/spin)

All tools fully documented with examples!

Future Directions

Upcoming features:

1. Enhanced Visualization

- ▶ Interactive MD trajectory viewer
- ▶ 3D ESP visualization tools
- ▶ Real-time training monitoring

2. Additional Models

- ▶ Graph neural networks (GNNs)
- ▶ Transformer architectures
- ▶ Ensemble methods

3. Advanced Features

- ▶ Active learning workflows
- ▶ Uncertainty quantification
- ▶ Transfer learning support

4. Integration

- ▶ More quantum chemistry packages
- ▶ Cloud training support
- ▶ HPC job submission

Acknowledgments

Thank you for using MMML!

MMML Development Team

Built with:

- ▶ JAX & Flax (neural networks)
- ▶ e3x (equivariant operations)
- ▶ ASE (atomic simulations)
- ▶ NumPy & SciPy (numerical computing)

Questions?

GitHub: [mmml](#)

Documentation: [docs/](#)